

Cloud Architecture Design Document

Overview

The **Cloud-Based Retail Inventory Management System** is built on **Amazon Web Services (AWS)** to provide a scalable, secure, and cost-effective solution for managing inventory, tracking sales, and generating real-time analytics. The system leverages a combination of AWS services, including **EC2**, **S3**, **RDS**, **Lambda**, **VPC**, and **CloudWatch**, to deliver a robust and efficient cloud infrastructure.

System Architecture

The system architecture is divided into the following components:

1. **Frontend:** A web-based dashboard hosted on an EC2 instance for real-time inventory and sales analytics.
 2. **Backend:** Python scripts and Lambda functions for automating inventory updates, sales reporting, and low-stock alerts.
 3. **Database:** A relational database (RDS) for storing inventory, sales, and user data.
 4. **Storage:** S3 buckets for storing inventory and sales data, as well as backups.
 5. **Networking:** A VPC with public and private subnets for secure communication between components.
 6. **Monitoring:** CloudWatch for monitoring system performance and setting up alarms.
-

AWS Services Used

1. Amazon EC2

- **Purpose:** Hosts the web application and Python scripts for inventory management.
- **Justification:** EC2 provides scalable compute capacity, allowing the system to handle varying loads. Auto Scaling groups ensure high availability and performance.

2. Amazon S3

- **Purpose:** Stores inventory and sales data, as well as backups.

- **Justification:** S3 is highly durable and scalable, making it ideal for storing large volumes of data. It also integrates seamlessly with other AWS services like Lambda and RDS.

3. Amazon RDS

- **Purpose:** Manages a relational database for inventory, sales, and user data.
- **Justification:** RDS was chosen over DynamoDB because the system requires complex queries and relationships between tables (e.g., products, inventory, sales). RDS supports SQL-based databases like MySQL and PostgreSQL, which are better suited for this use case.

4. AWS Lambda

- **Purpose:** Executes serverless functions for automating inventory updates, sales reporting, and low-stock alerts.
- **Justification:** Lambda eliminates the need to manage servers, reducing operational overhead. It is cost-effective for event-driven tasks and scales automatically.

5. Amazon VPC

- **Purpose:** Provides a secure and isolated network environment for the system.
- **Justification:** VPC allows for fine-grained control over network traffic, ensuring that sensitive data is protected. Public and private subnets are used to separate web-facing components from backend services.

6. Amazon CloudWatch

- **Purpose:** Monitors system performance, sets up alarms, and logs events.
- **Justification:** CloudWatch provides real-time insights into system health, enabling proactive troubleshooting and optimization.

Network Configuration

Virtual Private Cloud (VPC)

- **CIDR Block:** 10.0.0.0/16
- **Subnets:**
 - **Public Subnets:** 10.0.1.0/24 and 10.0.2.0/24 (for EC2 instances and web-facing components).
 - **Private Subnets:** 10.0.3.0/24 and 10.0.4.0/24 (for RDS and Lambda functions).
- **Internet Gateway:** Attached to the VPC for internet access.

- **Route Tables:** Configured to route traffic between subnets and the internet.

Security Groups

- **EC2 Security Group:** Allows SSH (port 22) and HTTP (port 80) traffic.
 - **RDS Security Group:** Allows inbound traffic from EC2 instances and Lambda functions.
 - **Lambda Security Group:** Allows outbound traffic to RDS and S3.
-

Security Measures

1. IAM Roles and Policies

- **EC2 Role:** Grants permissions to access S3, RDS, and CloudWatch.
- **Lambda Role:** Grants permissions to access S3, RDS, and CloudWatch.
- **RDS Role:** Grants permissions for automated backups and encryption.

2. Encryption

- **S3:** Server-Side Encryption (SSE-S3) for data at rest.
- **RDS:** Encryption enabled for the database.
- **Data in Transit:** SSL/TLS encryption for communication between components.

3. Security Groups

- Restrict access to EC2 instances and RDS databases to only trusted IP addresses.
 - Use private subnets for backend services to minimize exposure to the internet.
-

Scalability and Cost Optimization Strategies

1. Scalability

- **EC2 Auto Scaling:** Automatically adjusts the number of EC2 instances based on traffic.
- **RDS Read Replicas:** Improves database performance by offloading read queries.
- **Lambda:** Scales automatically to handle varying workloads.

2. Cost Optimization

- **Right-Sizing EC2 Instances:** Use instance types that match the workload requirements.
- **Reserved Instances:** Purchase reserved instances for long-term cost savings.

- **S3 Lifecycle Policies:** Move infrequently accessed data to S3 Glacier for cost savings.
 - **AWS Cost Explorer:** Monitor and optimize cloud expenses.
-

Conclusion

The **Cloud-Based Retail Inventory Management System** leverages AWS services to provide a scalable, secure, and cost-effective solution for retail businesses. The architecture is designed to handle varying workloads, ensure data security, and optimize costs. By using a combination of EC2, S3, RDS, Lambda, VPC, and CloudWatch, the system delivers real-time inventory tracking, automated stock replenishment, and detailed sales analytics.

Diagrams

Architecture Diagram

![alt text](architecture_diagram.png)

References

- [AWS Documentation](#)
- [AWS Well-Architected Framework](#)